An Introduction To Machine Learning[1]

Abdelbacet Mhamdi
✉ abdelbacet.mhamdi@bizerte.r-iset.tn

Dr.-Ing. in Electrical Engineering
Senior Lecturer at ISET Bizerte

"Computers are able to see, hear and learn.
Welcome to the future."

Dave Waters

"This is nothing. In a few years, that bot will move
so fast you'll need a strobe light to see it.
Sweet dreams…"

Elon Musk

"Machine intelligence is the last invention
that humanity will ever need to make."

Nick Bostrom

[1] Available @ https://github.com/a-mhamdi/mlpy/

Disclaimer

This document features some materials gathered from multiple online sources. Please note no copyright infringement is intended, and I do not own nor claim to own any of the original materials. They are used for educational purposes only. I have included links solely as a convenience to the reader. Some links within these slides may lead to other websites, including those operated and maintained by third parties. The presence of such a link does not imply a responsibility for the linked site or an endorsement of the linked site, its operator, or its contents.

Next...

Trends

"Numbers represent search interest relative to the highest point on the chart for the given region and time.

- A value of 100 is the peak popularity for the term;

- A value of 50 means that the term is half as popular;

- A score of 0 means there was not enough data for this term."

Global Data Traffic

The Internet in Real-Time
How Quickly Data is Generated

Twitter 330 Accounts Created / 171000 Tweets	**YouTube** 69420 Video Hours Watched / 60 Video Hours Uploaded	**LinkedIn** 5460 User Searches
Skype 694440 Minutes Used	**Instagram** 555570 Likes / 30000 Comments / 20820 Uploaded	

Google+ 173610 +1s	**WordPress** 690 Blog Posts	**Google** 138240 Searches / $48060 Ad Revenue
Reddit 30 Posts / 390 Comments / 6360 Votes	**Tumblr** 13890 Posts	**Pinterest** 7140 Pins

amazon 1530 Items Purchased / $70770 Money Spent	**foursquare** 1050 Check-Ins	**yelp** 15 Reviews
Email 102083340 Emails Sent	**Dropbox** 347220 Files Saved	**Snapchat** 173610 Stories Viewed / 243060 Messages Sent

Apple 19020 App Downloads	**Android** 37080 App Downloads	**Facebook** 1565880 Likes / 1649280 Posts / 180 GB of Data
WhatsApp 360 Accounts Created / 6597210 Messages Sent	**NETFLIX** 11580 Hours Watched	**PANDORA** 30570 Hours Streamed

Source: pennystocks.la/internet-in-real-time
Gif by: digitalinformationworld.com

By the way, in the 30 seconds you've been on this page, approximately 677220 GB of data was transferred over the internet.

Update on the internet in real time is available ▶ here .

Top Uses

Literature Review (1/1)

"The field of machine learning is concerned with the question of how to construct computer programs that automatically improve with experience."

Mitchell, T. (1997) *Machine Learning.* McGraw-Hill International Editions. McGraw-Hill.

Literature Review (2/1)

"Machine learning (ML) is a scientific discipline that concerns developing learning capabilities in computer systems. Machine learning is one of central areas of Artificial Intelligence (AI). It is an interdisciplinary area that combines results from statistics, logic, robotics, computer science, computational intelligence, pattern recognition, data mining, cognitive science, and more."

Wojtusiak, J. (2012) Machine learning. In *Encyclopedia of the Sciences of Learning*, pages 2082–2083. Springer US.

Literature Review (3/1)

"Machine learning is an evolving branch of computational algorithms that are designed to emulate human intelligence by learning from the surrounding environment. They are considered the working horse in the new era of the so-called big data. Techniques based on machine learning have been applied successfully in diverse fields ranging from pattern recognition, computer vision, spacecraft engineering, finance, entertainment, and computational biology to biomedical and medical applications. [...] The ability of machine learning algorithms to learn from current context and generalize into unseen tasks would allow improvements in both the safety and efficacy of radiotherapy practice leading to better outcomes."

El Naqa, I. and Murphy, M. J. (2015) *What Is Machine Learning?*, pages 3–11. Springer International Publishing.

Debrief

Arthur Samuel (1959)

Machine Learning: Field of study that gives computers the ability to learn without being explicitly programmed.

Tom Mitchell (1998)

Well-posed Learning Problem: A computer is said to learn from experience $\mathcal{E}$ with respect to some task $\mathcal{T}$ and some performance measure $\mathcal{P}$, if its performance on $\mathcal{T}$, as measured by $\mathcal{P}$, improves with experience $\mathcal{E}$.
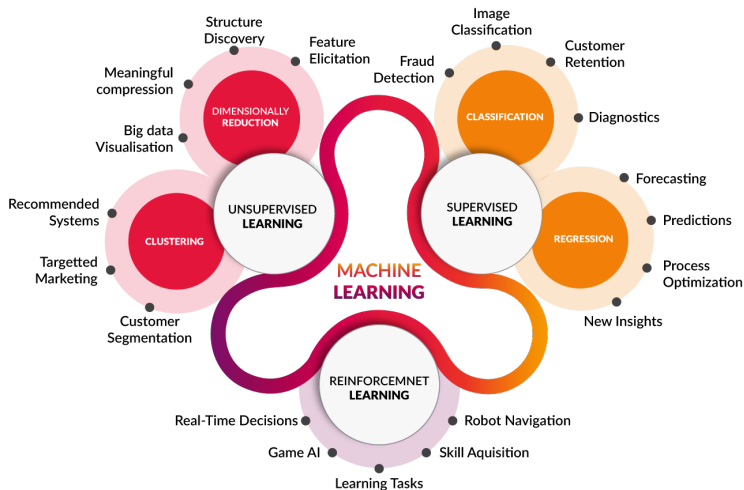
Debrief

Arthur Samuel (1959)

<u>Machine Learning:</u> Field of study that gives computers the ability to learn without being explicitly programmed.

Tom Mitchell (1998)

<u>Well-posed Learning Problem:</u> A computer is said to learn from experience $\mathcal{E}$ with respect to some task $\mathcal{T}$ and some performance measure $\mathcal{P}$, if its performance on $\mathcal{T}$, as measured by $\mathcal{P}$, improves with experience $\mathcal{E}$.

Task #1

Suppose your email program watches which emails you do or do not mark as spam, and based on that learns how to better filter spam. What is the task $\mathcal{T}$ in this setting?
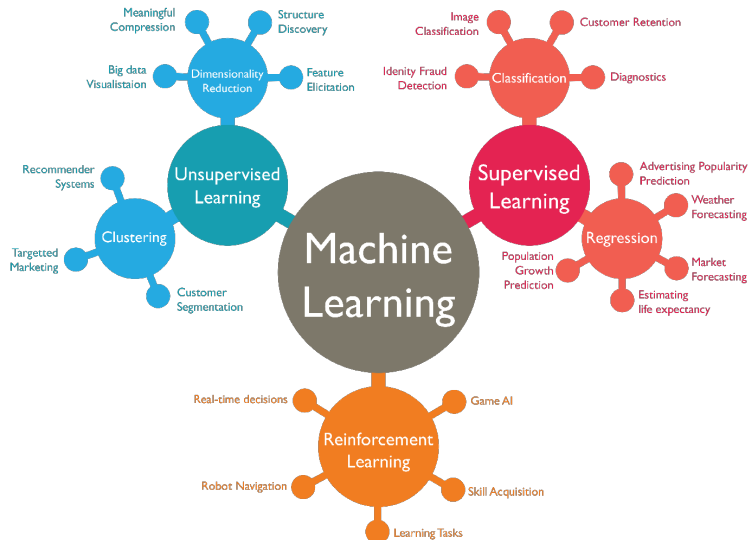
1. Classifying emails as spam or not spam;
2. Watching you label emails as spam or not spam;
3. The number (or fraction) of emails correctly classified as spam/not spam;
4. None of the above-this not a machine learning problem.

Debrief

Arthur Samuel (1959)

<u>Machine Learning:</u> Field of study that gives computers the ability to learn without being explicitly programmed.

Tom Mitchell (1998)

<u>Well-posed Learning Problem:</u> A computer is said to learn from experience $\mathcal{E}$ with respect to some task $\mathcal{T}$ and some performance measure $\mathcal{P}$, if its performance on $\mathcal{T}$, as measured by $\mathcal{P}$, improves with experience $\mathcal{E}$.

Task #1

Suppose your email program watches which emails you do or do not mark as spam, and based on that learns how to better filter spam. What is the task $\mathcal{T}$ in this setting?

1. Classifying emails as spam or not spam;
2. Watching you label emails as spam or not spam;
3. The number (or fraction) of emails correctly classified as spam/not spam;
4. None of the above-this not a machine learning problem.

Overall Methodology

1. Define the problem;
2. Gather dataset;
3. Choose measure of success;
4. Decide evaluation protocol;
5. Prepare the data;
6. Develop a model;
7. Iterate models.

MACHINE LEARNING

- **UNSUPERVISED LEARNING**
  - **DIMENSIONALLY REDUCTION**
    - Structure Discovery
    - Feature Elicitation
    - Meaningful compression
    - Big data Visualisation
  - **CLUSTERING**
    - Recommended Systems
    - Targetted Marketing
    - Customer Segmentation
- **SUPERVISED LEARNING**
  - **CLASSIFICATION**
    - Image Classification
    - Customer Retention
    - Fraud Detection
    - Diagnostics
  - **REGRESSION**
    - Forecasting
    - Predictions
    - Process Optimization
    - New Insights
- **REINFORCEMNET LEARNING**
  - Real-Time Decisions
  - Robot Navigation
  - Game AI
  - Skill Aquisition
  - Learning Tasks

https://www.cognub.com/index.php/cognitive-platform/

https://vitalflux.com/great-mind-maps-for-learning-machine-learning/

scikit-learn
algorithm cheat-sheet

https://scikit-learn.org/stable/tutorial/machine_learning_map/index.html

Regression | Classification | Clustering

A. Mhamdi   An Introduction To ML

Programming Language

python.org/

Development Environments

▲ $ docker compose up
▼ $ docker compose down
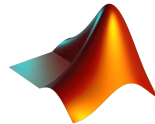
Required Packages

Valid only for...

python™

- A full list is available @ https://pypi.org/

Numpy

Matplotlib

Pandas

Scikit − learn

Keras

```
$ pip install virtualenv
$ virtualenv - -version
$ virtualenv «virtualenv_name»
$ source «virtualenv_name»/bin/activate  # ACTIVATE
$ deactivate  # DEACTIVATE
```

```
> pip install virtualenv
> virtualenv - -version
> virtualenv «virtualenv_name»
> «virtualenv_name»\Scripts\activate  .: ACTIVATE
> deactivate  .: DEACTIVATE
```

Source Control Management (SCM)

https://github.com/a-mhamdi/mlpy

Continuous Integration (CI)

https://hub.docker.com/r/abmhamdi/mlpy

Next...

Workflow in Machine Learning

Data Preprocessing
How?

Cleaning Identifying and correcting or removing inaccuracies and inconsistencies in the data.

Transformation Converting data from one format or structure to another.

Normalization Scaling the data so that it fits within a specific range. This is often done to make the data more amenable to certain operations or algorithms.

Data Preprocessing
Why?

▶ Raw data is often messy and may need to be cleaned and formatted before it can be used for machine learning.
(This may involve removing missing or invalid data, handling outliers, and encoding categorical variables.)

▶ Normalizing the data can help to scale the features so that they are on the same scale.
(This can be important for algorithms that use distance measures, as features on different scales can dominate the distance measure.)

▶ Preprocessing techniques such as feature selection and feature extraction can help to reduce the dimensionality of the data.
(This may improve the performance of the model and reduce the risk of overfitting.)

▶ Preprocessing techniques such as feature selection can help to identify the most important features in the data.
(This can make the model more interpretable and easier to understand.)

Data Preprocessing
Feature Scaling

Normalization *(MinMaxScaler)*	**Standardization** *(StandardScaler)*
$X \triangleq \dfrac{X - \min(X)}{\max(X) - \min(X)}$	$X \triangleq \dfrac{X - \mu}{\sigma}$
▲ No assumption on data distribution	▲ More recommended when following normal distribution

A. Mhamdi   An Introduction To ML

Data Preprocessing Template

```python
[ ]:  from sklearn.preprocessing import StandardScaler
```

```python
[ ]:  sc = StandardScaler()
      X_train = sc.fit_transform(X_train)
      X_test = sc.transform(X_test)
```

The notebook is available at https://github.com/a-mhamdi/mlpy/
→ *Codes* → *Python* → *data-preprocessing.ipynb*

This kind of supervised learning deals with labelled data. A subset of this data is used later to predict in continuous form. Regression problems involve tasks where the outputs form generally a set of real numbers. They oftenly follow linear formats.
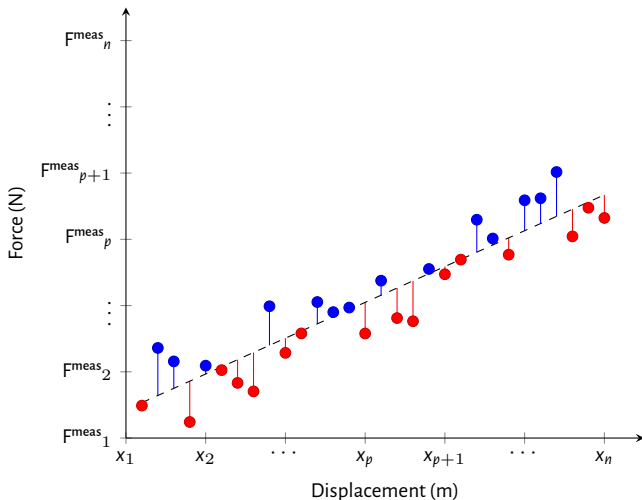
Consider the example of a spring. Our main goal is to determine the stiffness k of this spring, given some experimental data. The mathematical model *(Hooke's law)*:

$$F = kx \tag{1}$$

Restoring force is proportional to displacement.

Table: Measurements of couple (x_i, F^{meas}_i)

x_i	x_1	$\cdots$	x_p	$\cdots$	x_n
F^{meas}_i	F^{meas}_1	$\cdots$	F^{meas}_p	$\cdots$	F^{meas}_n

$$
\begin{aligned}
F^{meas}_i &= F_i + \varepsilon_i \\
&= kx_i + \varepsilon_i, \tag{2}
\end{aligned}
$$

where F_i denotes the unknown real value of the force applied to the spring. In order to estimate the stiffness value k, we can consider the quadratic criterion:

$$
\begin{aligned}
\mathcal{J} &= \sum_{i=1}^{n} \varepsilon_i^2 \\
&= \sum_{i=1}^{n} \left(F^{meas}_i - kx_i \right)^2
\end{aligned}
$$

$$\frac{\partial \mathcal{J}}{\partial k} \;=\; 0 \tag{3}$$

$$2 \sum_{i=1}^{n} \left(F^{\text{meas}}{}_i - kx_i \right) \sum_{i=1}^{n} \frac{\partial \left(F^{\text{meas}}{}_i - kx_i \right)}{\partial k} \;=\; 0$$

$$\sum_{i=1}^{n} \left(F^{\text{meas}}{}_i - kx_i \right) \sum_{i=1}^{n} x_i \;=\; 0$$

$$\sum_{i=1}^{n} F^{\text{meas}}{}_i \, x_i \;=\; k \sum_{i=1}^{n} x_i^2 \quad \Longleftrightarrow \quad \boxed{\hat{k} \;=\; \frac{\sum_{i=1}^{n} F^{\text{meas}}{}_i \, x_i}{\sum_{i=1}^{n} x_i^2}}$$

Simple Linear Regression

Training the Simple Linear Regression model on the Training set

```
[ ]:  from sklearn.linear_model import LinearRegression
```

```
[ ]:  lr = LinearRegression()
      lr.fit(X_train, y_train)
```

```
[ ]:  LinearRegression()
```

Predicting the Test set results

```
[ ]:  y_pred = lr.predict(X_test)
```

The notebook is available at https://github.com/a-mhamdi/mlpy/
→ *Codes* → *simple-linear-regression.ipynb*

Multiple Linear Regression

Training the multiple linear regression model on the training set

```
[ ]:  from sklearn.linear_model import LinearRegression
```

```
[ ]:  lr = LinearRegression()
      lr.fit(X_train, y_train)
```

```
[ ]:  LinearRegression()
```

Making predictions using the X test set and comparison

```
[ ]:  y_pred = lr.predict(X_test)
```

The notebook is available at https://github.com/a-mhamdi/mlpy/
→ *Codes* → *multiple-linear-regression.ipynb*

Consider the following multivariate equation:

$$y = \theta_1 x_{(1)} + \theta_2 x_{(2)} + \cdots + \theta_m x_{(m)} \tag{4}$$

For a particular single measurement, eq. (**??**) can be updated as

$$y_k = \theta_1 x_{(1, k)} + \theta_2 x_{(2, k)} + \cdots + \theta_m x_{(m, k)} + \varepsilon_k \tag{5}$$

We denote hereafter by θ the vector $\begin{bmatrix} \theta_1 \\ \theta_2 \\ \vdots \\ \theta_n \end{bmatrix}$. The function y_k becomes:

$$y_k = \underbrace{\left[x_{(1, k)}, x_{(2, k)}, \cdots, x_{(m, k)} \right]}_{x_k^T} \theta + \varepsilon_k \tag{6}$$

We assume that we have n measurements for y. Then we can transform the previous equation into

$$y = X\theta + \varepsilon, \tag{7}$$

where $y^T = [y_1, y_2, \cdots, y_n]$, $X = \begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_n^T \end{bmatrix}$, and $\varepsilon^T = [\varepsilon_1, \varepsilon_2, \cdots, \varepsilon_n]$.

The regressor is given by

$$\hat{\theta} = \left(X^T X\right)^{-1} X^T y$$

$X^T X$ **is not invertible (singular/degenerate)**

▼ Redundant Features

Some features are linearly dependant, i.e, $\exists$ some $x_p \propto$ some x_l.

▼ Too many features

Fewer observations compared to the number of features, i.e, $m \geq n$.

▲ Delete some features
▲ Add extra observations

$$\left\{ \text{GRADIENT DESCENT} \right\}$$

The linear regression is given by:

$$y_k = \underbrace{\underbrace{\left[x_{(1,k)}, x_{(2,k)}, \cdots, x_{(m,k)} \right]}_{x_k^T} \theta + \varepsilon_k}_{h_\theta(x_k)} \tag{8}$$

$$\theta_0 \triangleq \theta_0 + \alpha \frac{1}{n} \sum_{k=1}^{n} (y_k - h_\theta(x_k)) x_{(0,k)}$$

$$\theta_1 \triangleq \theta_1 + \alpha \frac{1}{n} \sum_{k=1}^{n} (y_k - h_\theta(x_k)) x_{(1,k)}$$

$$\vdots$$

$$\theta_m \triangleq \theta_m + \alpha \frac{1}{n} \sum_{k=1}^{n} (y_k - h_\theta(x_k)) x_{(m,k)}$$

The parameter α is the learning rate.

Polynomial Regression

```
[ ]:  from sklearn.preprocessing import PolynomialFeatures
```

```
[ ]:  poly_reg = PolynomialFeatures(degree=4)
      X_poly = poly_reg.fit_transform(X)
      print(X_poly[:5])
      lr_2 = LinearRegression()
      lr_2.fit(X_poly, y)
```

The notebook is available at https://github.com/a-mhamdi/mlpy/
→ *Codes* → *Python* → *polynomial-regression.ipynb*

Task #2

The yield y of a chemical process is a random variable whose value is considered to be a linear function of the temperature x. The following data of corresponding values of x and y is found:

Temperature in °C (x)	0	25	50	75	100
Yield in grams (y)	14	38	54	76	95

The linear regression model $y = \theta_0 + \theta_1 x$ is used. Determine the values of θ_0, θ_1.

1. Using normal equation,

2. Using gradient descent for 5 iterations.

$$
y = \begin{bmatrix} 14 \\ 38 \\ 54 \\ 76 \\ 95 \end{bmatrix} \quad \text{and} \quad X = \begin{bmatrix} 1 & 0 \\ 1 & 25 \\ 1 & 50 \\ 1 & 75 \\ 1 & 100 \end{bmatrix} \implies X^T X = \begin{bmatrix} 5 & 250 \\ 250 & 18750 \end{bmatrix}
$$

$$
\hat{\theta} = \begin{bmatrix} \hat{\theta}_0 \\ \hat{\theta}_1 \end{bmatrix} = \begin{bmatrix} 15.4 \\ 0.8 \end{bmatrix}
$$

```python
import matplotlib.pyplot as plt
import numpy as np

X = np.array([[1,0], [1,25], [1,50], [1,75], [1,100]], dtype=np.float32)
y = np.array([[14], [38], [54], [76], [95]])

# NORMAL EQUATION
XtX = X.T.dot(X)
invXtX = np.linalg.inv(XtX)
t_ne = invXtX.dot(np.matmul(X.T, y))
print(t_ne)
X[:,1]= (X[:,1]-X[:,1].min())/X[:,1].max()
y = (y-y.min())/y.max()

# GRADIENT DESCENT
t_gd = np.array([[1], [1]])
alpha = .1
vect = np.zeros(shape=(2, 1001))
vect[:,0] = t_gd[[0,1],[0]]
lost = []
for k in range(1000):
    eps = y-np.matmul(X, t_gd)
    lost.append(1/(2*len(y))*eps.T.dot(eps)[[0],[0]][0])
    t_gd = t_gd+alpha*1/len(y)*np.matmul((eps).T, X).T
```

```
    vect[:,k+1] = t_gd[[0,1],[0]]

print(vect[:, -1])
plt.plot(vect[0, :], label=r'$\hat{\theta}_0$')
plt.plot(vect[1, :], label=r'$\hat{\theta}_1$')
plt.legend(); plt.grid(); plt.show()

plt.plot(lost); plt.grid(); plt.show()
```

F1-Score, **Accuracy**, **Recall** and **Precision** are calculated as follow:

$$f1 - \text{score} \;=\; \frac{2}{\dfrac{1}{\text{Recall}} + \dfrac{1}{\text{Precision}}}$$

$f1 -$ score denotes the *Harmonic Mean of Recall & Precision*

$$\text{Accuracy} \;=\; \frac{\text{TP} + \text{TN}}{\text{TP} + \text{FP} + \text{TN} + \text{FN}}$$

It denotes the ratio of how much we got right over all cases. Recall, on the other hand, designates the ratio of how much positives do we got right over all actual positive cases.

$$\text{Recall} \;=\; \frac{\text{TP}}{\text{TP} + \text{FN}}$$

Precision, at last, is how much positives we got right over all positive predictions. It is given by:

$$\text{Precision} \;=\; \frac{\text{TP}}{\text{TP} + \text{FP}}$$

Logistic or S-shaped function σ

$$\sigma(x) \;=\; \frac{1}{1 + \mathrm{e}^{-x}}$$

▲ σ squashes range of distance from $]-\infty,\ +\infty[$ to $[0,\ 1]$

▲ σ is differentiable and easy to compute: $\dot{\sigma} = \sigma \times (1 - \sigma)$

Decision boundary

$$y = \sigma\left(\theta_1 x_{(1)} + \theta_2 x_{(2)} + \cdots + \theta_m x_{(m)}\right)$$

$$y = \frac{1}{1 + e^{-\theta^T x}}$$

Hypothesis:

$$\text{Considering} \quad h_\theta\left(x\right) = \frac{1}{1 + e^{-\theta^T x}} \quad \text{yields} \quad h_\theta\left(x_k\right) = \frac{1}{1 + e^{-\theta^T x_k}}$$

$$\boxed{\textsc{Gradient Descent}}$$

$$\vdots$$

$$\boxed{\theta_m \triangleq \theta_m + \alpha \frac{1}{n} \sum_{k=1}^{n} \left(y_k - h_\theta\left(x_k\right)\right) x_{(m,\,k)}}$$

Logistic Regression

Training the logistic regressor

```
[ ]: from sklearn.linear_model import LogisticRegression
```

```
[ ]: clf = LogisticRegression(random_state=123)
     clf.fit(X_train, y_train)
```

```
[ ]: LogisticRegression(random_state=123)
```

Predicting the test set results

```
[ ]: y_pred = clf.predict(X_test)
```

The notebook is available at https://github.com/a-mhamdi/mlpy/
→ *Codes* → *Python* → *logistic-regression.ipynb*

k-Nearest Neighbors (1/1)

▶ Evelyn Fix and Joseph Hodges, 1951   ▶ Thomas Cover, 1966

Algorithm Summary Construction

1: **procedure** How does k-NN work? (Finding Nearest Neighbors)
   Input: A query point;
   Output: Assign a class label to that point.
2:     Define how many neighbors will be checked to classify the specific query point;
3:     Compute the distance $d(x; y)$ of the query point to other data points;
4:     Count the number of the data points in each category;
5:     Assign the query point to the class with most frequent neighbors.
6: **end procedure**

k-Nearest Neighbors (2/1)

Minkowski distance
$$d(x; y) = \left(\sum_{i=1}^{n} |y_i - x_i|^p \right)^{1/p}$$

Manhattan distance (p=1)
$$d(x; y) = \sum_{i=1}^{n} |y_i - x_i|$$

Euclidean distance (p=2)
$$d(x; y) = \sqrt{\sum_{i=1}^{n} (y_i - x_i)^2}$$

Task #3

Let be the following coordinate points:

$A(1, 6)$; $B(2, 6)$; $C(3, 1)$; $D(4, 2)$; $E(6, 0)$; $F(7, 5)$; $G(7, 3)$; $H(10, 3)$; $I(-4, -1)$

Using the Euclidean distance, what are the two closest neighbors of point $P(5, 5)$?

$d(A; P) = \sqrt{17} \approx 4.12$ $d(B; P) = \sqrt{10} \approx 3.16$ $d(C; P) = \sqrt{20} \approx 4.47$

$d(D; P) = \sqrt{10} \approx 3.16$ $d(E; P) = \sqrt{26} \approx 5.1$ $d(F; P) = \sqrt{4} = 2$

$d(G; P) = \sqrt{8} \approx 2.83$ $d(H; P) = \sqrt{29} \approx 5.38$ $d(I; P) = \sqrt{117} \approx 10.82$

k-Nearest Neighbors (3/1)

```python
from math import sqrt
def dds(a, b): # `a` and `b` are coordinates of some point
    d_squared = (a-5)**2+(b-5)**2
    return (d_squared, sqrt(d_squared))

dds(1, 6) # Point `A`
dds(2, 6) # Point `B`
```

k-Nearest Neighbors (4/1)

Task #4[a]

[a]From Prof. Winston's book

We try to predict the color of a fruit according to its width (w) and height (h). The following training data is available:

Fruit	F_1	F_2	F_3	F_4	F_5	F_6	F_7	F_8
w	2	5	2	6	1	4	2	6
h	6	6	5	5	2	2	1	1
Color	Red	Yellow	Orange	Purple	Red	Blue	Violet	Green

The goal here is to study the influence of neighbors on the color property of a fruit. Let U be the new fruit of width $w = 1$ and height $h = 4$

1. What is its color if we consider 1 neighbor?
2. What is its color if we consider 3 neighbors?
3. Rather than majority voting, we would like to consider the vote of neighbors weighted by the distance. Each neighbor votes according to a weight inversely proportional to the square of its distance: $\frac{1}{d^2}$. We take 3 neighbors, what is the color of U? Compare your results to those in question 2.

k-Nearest Neighbors (5/1)

$$d(U; F_1) = \sqrt{5} \approx 2.24 \qquad d(U; F_2) = \sqrt{20} \approx 4.47 \qquad d(U; F_3) = \sqrt{2} \approx 1.41$$

$$d(U; F_4) = \sqrt{26} \approx 5.1 \qquad d(U; F_5) = \sqrt{4} = 2 \qquad d(U; F_6) = \sqrt{13} \approx 3.6$$

$$d(U; F_7) = \sqrt{10} \approx 3.16 \qquad d(U; F_8) = \sqrt{34} \approx 5.83$$

1. Color of U is Orange because $d(U; F_3)$ is the smallest.
2. Color of U is Red: F_1 and F_5 (+2 to Red class), F_3 (+1 to Orange class)
3. Color of U is Orange

$$\mathcal{S}(\text{Red}) = \frac{1}{d^2(U; F_1)} + \frac{1}{d^2(U; F_5)} = 0.45 \qquad \mathcal{S}(\text{Orange}) = \frac{1}{d^2(U; F_3)} = 0.5$$

```python
from math import sqrt
def dds(w, h): # `w` and `h` are width and height of some fruit
    d_squared = (w-1)**2+(h-4)**2
    return (d_squared, sqrt(d_squared))

dds(2, 6) # Fruit `F_1`
dds(5, 6) # Fruit `F_2`
```

k-NN

Importing the classifier

```
[ ]: from sklearn.neighbors import KNeighborsClassifier
```

Training the K-NN model on the training set

```
[ ]: clf = KNeighborsClassifier(n_neighbors, metric, p)
     clf.fit(X_train, y_train)
```

Predicting the test set results

```
[ ]: y_pred = clf.predict(X_test)
```

The notebook is available at https://github.com/a-mhamdi/mlpy/
→ *Codes* → *Python* → *k-nearest-neighbors.ipynb*

Rule of Thumb to Choose k

k is even if the number of classes is odd

k is odd if the number of classes is even

k is an important hyperparameter that can affect the performance of the model.

1. Larger values of k will result in a smoother decision boundary, which can lead to a more generalized model.
2. Smaller values of k will result in a more complex decision boundary, which can lead to a model that is more prone to overfitting.
3. The optimal value of K may depend on the specific dataset and the characteristics of the data.

Outroduction

Method		Pros		Cons
Logistic Regression	▲	Probabilistic	▼	Almost linearly separable data
K-NN	▲	Simple	▼	Number of neighbors k
	▲	Fast	▼	Detecting outliers[2]
	▲	Efficient		

---

[2]Points that differ significantly from the rest of the data points.

A. Mhamdi An Introduction To ML

Next...

K-Means Clustering (1/1)

The algorithm K-**Means** allows to display regularities or patterns in unlabeled data.

▶ The term 'means' refers to averaging the data when computing each centroid;

▶ A centroid is the arithmetic mean of all the data points belonging to a particular cluster.

This technique identifies a certain number of centroids within a data set. The algorithm then allocates every data point to the nearest cluster as it attempts to keep the clusters as small as possible. At the same time, K-Means attempts to keep the other clusters as different as possible.

K-Means Clustering (2/1)

K-Means Clustering (3/1)

Algorithm Summary Construction

1: **procedure** HOW DOES *K*-MEANS WORK? (Discovering similarities)
   Input: Unlabeled data sets;
   Output: Grouping into clusters.
2:    Define how many clusters will be used to group the data sets;
3:    Initialize all the coordinates of the *k* cluster centers
4:    **repeat**
5:        Assign each point to its nearest cluster;
6:        Update the centroids coordinates;
7:    **until** No changes to the centers of the clusters
8:    Assign new cases to one of the clusters
9: **end procedure**

A. MHAMDI   An Introduction To ML

Task #5[a]

[a]From 'Machine Learning' course on 'Coursera'

Of the following examples, which would you address using an underlined unsupervised learning algorithms? *(Check all that apply.)*

1. Given email labeled as spam/not spam, learn a spam filter
2. Given a set of news articles found on the web, group them into set of articles about the same story
3. Given a database of customer data, automatically discover market segments and group customers into different market segments
4. Given a dataset of patients diagnosed as either having diabetes or not, learn to classify new patients as having diabetes or not.

Task #5[a]

[a]From 'Machine Learning' course on 'Coursera'

Of the following examples, which would you address using an <u>unsupervised learning</u> algorithms? *(Check all that apply.)*

1. Given email labeled as spam/not spam, learn a spam filter
2. Given a set of news articles found on the web, group them into set of articles about the same story
3. Given a database of customer data, automatically discover market segments and group customers into different market segments
4. Given a dataset of patients diagnosed as either having diabetes or not, learn to classify new patients as having diabetes or not.

Task #6[a]

[a]Credit: Shokoufeh Mirzaei, PhD

Use *K*-Means algorithm to cluster the following eight points intro three clusters:

$$A(2, 10); \; B(2, 5); C(8, 4); D(5, 8); E(7, 5); F(6, 4); G(1, 2) \text{ and } H(4, 9).$$

- Initial cluster centers are: $\alpha(2, 10); \beta(5, 8)$ and $\gamma(1, 2)$
- The distance between two points: $M(x_m, y_m)$ and $N(x_n, y_n)$ is defined as

$$d(M; N) \; = \; |x_m - x_n| + |y_m - y_n|$$

Task #6[a]

[a]Credit: Shokoufeh Mirzaei, PhD

Use *K*-Means algorithm to cluster the following eight points intro three clusters:

$A(2, 10); B(2, 5); C(8, 4); D(5, 8); E(7, 5); F(6, 4); G(1, 2)$ and $H(4, 9)$.

- ▶ Initial cluster centers are: $\alpha(2, 10); \beta(5, 8)$ and $\gamma(1, 2)$
- ▶ The distance between two points: $M(x_m, y_m)$ and $N(x_n, y_n)$ is defined as

$$d(M; N) = |x_m - x_n| + |y_m - y_n|$$

A. Mhamdi   An Introduction To ML

Task #6[a]

[a]Credit: Shokoufeh Mirzaei, PhD

Use *K*-Means algorithm to cluster the following eight points intro three clusters:

$$A(2, 10); \ B(2, 5); \ C(8, 4); \ D(5, 8); \ E(7, 5); \ F(6, 4); \ G(1, 2) \text{ and } H(4, 9).$$

- ▶ Initial cluster centers are: $\alpha(2, 10); \ \beta(5, 8)$ and $\gamma(1, 2)$
- ▶ The distance between two points: $M(x_m, y_m)$ and $N(x_n, y_n)$ is defined as

$$d(M; N) = |x_m - x_n| + |y_m - y_n|$$

Point	$\alpha(2, 10)$	$\beta(5, 8)$	$\gamma(1, 2)$	#
$A(2, 10)$	0	5	9	1
$B(2, 5)$	5	6	4	3
$C(8, 4)$	12	7	9	2
$D(5, 8)$	5	0	10	2
$E(7, 5)$	10	5	9	2
$F(6, 4)$	10	5	7	2
$G(1, 2)$	9	10	0	3
$H(4, 9)$	3	2	10	2

A. Mhamdi   An Introduction To ML

Task #6[a]

[a]Credit: Shokoufeh Mirzaei, PhD

Use *K*-Means algorithm to cluster the following eight points intro three clusters:

$$A(2, 10); \ B(2, 5); C(8, 4); D(5, 8); E(7, 5); F(6, 4); G(1, 2) \text{ and } H(4, 9).$$

► Initial cluster centers are: $\alpha(2, 10); \beta(5, 8)$ and $\gamma(1, 2)$
► The distance between two points: $M(x_m, y_m)$ and $N(x_n, y_n)$ is defined as

$$d(M; N) \ = \ |x_m - x_n| + |y_m - y_n|$$

Point	$\alpha(2, 10)$	$\beta(5, 8)$	$\gamma(1, 2)$	#
$A(2, 10)$	0	5	9	1
$B(2, 5)$	5	6	4	3
$C(8, 4)$	12	7	9	2
$D(5, 8)$	5	0	10	2
$E(7, 5)$	10	5	9	2
$F(6, 4)$	10	5	7	2
$G(1, 2)$	9	10	0	3
$H(4, 9)$	3	2	10	2
	$\alpha(2, 10)$	$\beta(6, 6)$	$\gamma(1.5, 3.5)$	

A. Mhamdi An Introduction To ML

Task #6[a]

[a]Credit: Shokoufeh Mirzaei, PhD

Use *K*-Means algorithm to cluster the following eight points intro three clusters:

$$A(2, 10); B(2, 5); C(8, 4); D(5, 8); E(7, 5); F(6, 4); G(1, 2) \text{ and } H(4, 9).$$

▶ Initial cluster centers are: $\alpha(2, 10); \beta(5, 8)$ and $\gamma(1, 2)$

▶ The distance between two points: $M(x_m, y_m)$ and $N(x_n, y_n)$ is defined as

$$d(M; N) = |x_m - x_n| + |y_m - y_n|$$

Point	$\alpha(2, 10)$	$\beta(6, 6)$	$\gamma(1.5, 3.5)$	#
$A(2, 10)$	0	8	7	1
$B(2, 5)$	5	5	2	3
$C(8, 4)$	12	4	7	2
$D(5, 8)$	5	3	8	2
$E(7, 5)$	10	2	7	2
$F(6, 4)$	10	2	5	2
$G(1, 2)$	9	9	2	3
$H(4, 9)$	3	5	8	1

Task #6[a]

[a]Credit: Shokoufeh Mirzaei, PhD

Use *K*-Means algorithm to cluster the following eight points intro three clusters:

$$A(2, 10); B(2, 5); C(8, 4); D(5, 8); E(7, 5); F(6, 4); G(1, 2) \text{ and } H(4, 9).$$

► Initial cluster centers are: $\alpha(2, 10); \beta(5, 8)$ and $\gamma(1, 2)$

► The distance between two points: $M(x_m, y_m)$ and $N(x_n, y_n)$ is defined as

$$d(M; N) = |x_m - x_n| + |y_m - y_n|$$

Point	$\alpha(2, 10)$	$\beta(6, 6)$	$\gamma(1.5, 3.5)$	#
$A(2, 10)$	0	8	7	1
$B(2, 5)$	5	5	2	3
$C(8, 4)$	12	4	7	2
$D(5, 8)$	5	3	8	2
$E(7, 5)$	10	2	7	2
$F(6, 4)$	10	2	5	2
$G(1, 2)$	9	9	2	3
$H(4, 9)$	3	5	8	1
$\boxed{\alpha(3, 9.5)}$	$\boxed{\beta(6.5, 5.25)}$	$\boxed{\gamma(1.5, 3.5)}$		

A. Mhamdi   An Introduction To ML

> ## Task #6[a]
>
> [a]Credit: Shokoufeh Mirzaei, PhD
>
> Use K-Means algorithm to cluster the following eight points intro three clusters:
>
> $A(2, 10); B(2, 5); C(8, 4); D(5, 8); E(7, 5); F(6, 4); G(1, 2)$ and $H(4, 9)$.
>
> ▶ Initial cluster centers are: $\alpha(2, 10); \beta(5, 8)$ and $\gamma(1, 2)$
> ▶ The distance between two points: $M(x_m, y_m)$ and $N(x_n, y_n)$ is defined as
>
> $$d(M; N) = |x_m - x_n| + |y_m - y_n|$$

Point	$\alpha(3, 9.5)$	$\beta(6.5, 5.25)$	$\gamma(1.5, 3.5)$	#
$A(2, 10)$	1.5	9.25	7	1
$B(2, 5)$	5.5	4.75	2	3
$C(8, 4)$	10.5	2.75	7	2
$D(5, 8)$	3.5	4.25	8	1
$E(7, 5)$	8.5	0.75	7	2
$F(6, 4)$	8.5	1.75	5	2
$G(1, 2)$	9.5	8.75	2	3
$H(4, 9)$	1.5	6.25	8	1

A. MHAMDI An Introduction To ML

Task #6[a]

[a]Credit: Shokoufeh Mirzaei, PhD

Use K-Means algorithm to cluster the following eight points intro three clusters:

$A(2, 10)$; $B(2, 5)$; $C(8, 4)$; $D(5, 8)$; $E(7, 5)$; $F(6, 4)$; $G(1, 2)$ and $H(4, 9)$.

▶ Initial cluster centers are: $\alpha(2, 10)$; $\beta(5, 8)$ and $\gamma(1, 2)$
▶ The distance between two points: $M(x_m, y_m)$ and $N(x_n, y_n)$ is defined as

$$d(M; N) = |x_m - x_n| + |y_m - y_n|$$

Point	$\alpha(3, 9.5)$	$\beta(6.5, 5.25)$	$\gamma(1.5, 3.5)$	#
$A(2, 10)$	1.5	9.25	7	1
$B(2, 5)$	5.5	4.75	2	3
$C(8, 4)$	10.5	2.75	7	2
$D(5, 8)$	3.5	4.25	8	1
$E(7, 5)$	8.5	0.75	7	2
$F(6, 4)$	8.5	1.75	5	2
$G(1, 2)$	9.5	8.75	2	3
$H(4, 9)$	1.5	6.25	8	1
	$\boxed{\alpha(3.67, 9)}$	$\boxed{\beta(7, 4.3)}$	$\boxed{\gamma(1.5, 3.5)}$	

Task #6[a]

[a]Credit: Shokoufeh Mirzaei, PhD

Use *K*-Means algorithm to cluster the following eight points intro three clusters:

$$A(2, 10); \ B(2, 5); \ C(8, 4); \ D(5, 8); \ E(7, 5); \ F(6, 4); \ G(1, 2) \text{ and } H(4, 9).$$

▶ Initial cluster centers are: $\alpha(2, 10); \beta(5, 8)$ and $\gamma(1, 2)$
▶ The distance between two points: $M(x_m, y_m)$ and $N(x_n, y_n)$ is defined as

$$d(M; N) \ = \ |x_m - x_n| + |y_m - y_n|$$

Point	$\alpha(3.67, 9)$	$\beta(7, 4.3)$	$\gamma(1.5, 3.5)$	#
$A(2, 10)$	2.67	10.7	7	1
$B(2, 5)$	5.67	5.7	2	3
$C(8, 4)$	9.33	1.3	7	2
$D(5, 8)$	2.33	5.7	8	1
$E(7, 5)$	7.33	0.7	7	2
$F(6, 4)$	7.33	1.3	5	2
$G(1, 2)$	9.67	8.3	2	3
$H(4, 9)$	0.33	7.7	8	1

A. Mhamdi   An Introduction To ML

Task #6[a]

[a]Credit: Shokoufeh Mirzaei, PhD

Use K-Means algorithm to cluster the following eight points intro three clusters:

$A(2, 10); B(2, 5); C(8, 4); D(5, 8); E(7, 5); F(6, 4); G(1, 2)$ and $H(4, 9)$.

▶ Initial cluster centers are: $\alpha(2, 10); \beta(5, 8)$ and $\gamma(1, 2)$
▶ The distance between two points: $M(x_m, y_m)$ and $N(x_n, y_n)$ is defined as

$$d(M; N) = |x_m - x_n| + |y_m - y_n|$$

Point	$\alpha(3.67, 9)$	$\beta(7, 4.3)$	$\gamma(1.5, 3.5)$	#
$A(2, 10)$	2.67	10.7	7	1
$B(2, 5)$	5.67	5.7	2	3
$C(8, 4)$	9.33	1.3	7	2
$D(5, 8)$	2.33	5.7	8	1
$E(7, 5)$	7.33	0.7	7	2
$F(6, 4)$	7.33	1.3	5	2
$G(1, 2)$	9.67	8.3	2	3
$H(4, 9)$	0.33	7.7	8	1
$\alpha(3.67, 9)$	$\beta(7, 4.3)$	$\gamma(1.5, 3.5)$		

A. Mhamdi   An Introduction To ML

K-Means

Import `KMeans` **class**

```
[ ]:   from sklearn.cluster import KMeans
```

Training the K-Means model on the dataset

```
[ ]:   kmeans = KMeans(n_clusters, init, random_state)
       y_pred = kmeans.fit_predict(X)
```

The notebook is available at https://github.com/a-mhamdi/mlpy/
→ *Codes* → *Python* → *k-means-clustering.ipynb*

Next...

Fundamental unit of a neural network (1/1)


Artificial neuron

Biological neuron

https://id.wikipedia.org/wiki/Sel_saraf

Fundamental unit of a neural network (2/1)

Task #7

Compute the output of the following neuron.

$$y = \text{sign}(1 \times -1 + 2.6 \times 0.7 + 0.7 \times 2.5 - 2 \times 0.15 + 0.8 \times 1.3) = 1$$

Fundamental unit of a neural network (3/1)

Multilayer Perceptron (MLP)

Task #8

For the above structure, determine how many parameters are to be adjusted.

$$\text{Params \#} = 5 \times 3 + 5 + 2 \times 5 + 2 = 32$$

✓ Design a structure

✓ Specify a loss function to minimize

✓ Optimize using gradient descent

    ① Feedforward propagation *(matrix multiplication and point-wise activation)*

    ② Back propagation *(multivariate chain rule)*

    ③ Update the weights accordingly

    A. MHAMDI   An Introduction To ML

Tinker with a neural network

https://playground.tensorflow.org/

ANN

```python
from keras.models import Sequential
from keras.layers import Dense
```

```python
clf = Sequential()
ndim = X_train.shape[1]
clf.add(Dense(units=8, activation="relu", input_dim=ndim))
clf.add(Dense(units=4, activation="relu"))
clf.add(Dense(units=4, activation="relu"))
clf.add(Dense(units=1, activation="sigmoid"))
```

```python
clf.compile(optimizer="adam", loss="binary_crossentropy",
            metrics=["accuracy"])
```

```python
clf.fit(X_train, y_train, batch_size=16, epochs=32);
```

The notebook is available at https://github.com/a-mhamdi/mlpy/
→ Codes → Python → artificial-neural-network.ipynb

List of available optimizers

Here is a list of some common optimizers for artificial neural networks:

$$\Delta \hat{\mathcal{W}} \triangleq \mathcal{F}\left(\nabla \underbrace{\mathcal{J}\left(\hat{\mathcal{W}}\right)}_{\text{Loss Function}} \right) \equiv \boxed{\hat{\mathcal{W}} \triangleq \hat{\mathcal{W}} + \mathcal{F}\left(\nabla \mathcal{J}\left(\hat{\mathcal{W}}\right)\right)} \quad \nabla \mathcal{J}\left(\hat{\mathcal{W}}\right) = \begin{bmatrix} \dfrac{\partial \mathcal{J}}{\partial \hat{w}_0} \\ \vdots \\ \dfrac{\partial \mathcal{J}}{\partial \hat{w}_n} \end{bmatrix}$$

SGD

SGD+MOMENTUM

ADAGRAD

WINDOWGRAD

ADADELTA

NESTEROV

Effect of optimizer on loss values

https://cs.stanford.edu/people/karpathy/convnetjs/demo/trainers.html

Effect of optimizer on testing accuracy values

https://cs.stanford.edu/people/karpathy/convnetjs/demo/trainers.html

Next...

An Introduction To ML

MCQ (1/1)

1 What is Machine Learning (ML)?

    ✗ The selective acquisition of knowledge through the use of computer programs

    ✗ The selective acquisition of knowledge through the use of manual programs

    ✓ The autonomous acquisition of knowledge through the use of computer programs

    ✗ The autonomous acquisition of knowledge through the use of manual programs

2 Successful applications of ML

    ✗ Learning to recognize spoken words

    ✗ Learning to drive an autonomous vehicle

    ✗ Learning to classify new astronomical structures

    ✗ Learning to play world-class backgammon

    ✓ All of the above

3 Features of Machine Learning are ...

    ✓ Automation

    ✓ Improved customer experience

    ✓ Business intelligence

4 ... is the machine learning algorithm that can be used with labeled data.

    ✓ Regression algorithm

    ✗ Clustering algorithm

    ✗ Association algorithm

MCQ (2/1)

5 Replace missing values with mean/median/mode helps to handle missing or corrupted data in a dataset. True/False?
- ✓ True
- ✗ False

6 Which among the following algorithms are used in Machine learning?
- ✓ Naive Bayes
- ✓ Support Vector Machines
- ✓ k-Nearest Neighbors

7 The term machine learning was coined by ...
- ✗ James Gosling
- ✓ Arthur Samuel
- ✗ Guido van Rossum
- ✗ None of the above

8 The Real-world machine learning use cases are
- ✓ Digital assistants
- ✓ Chatbots
- ✓ Fraud detection

MCQ (3/1)

9 Machine learning approaches can be traditionally categorized into ... categories.

- ✓ 3
- ✗ 4
- ✗ 7
- ✗ 9

10 ... is a part of machine learning that works with neural networks.

- ✗ Artificial intelligence
- ✓ Deep learning
- ✗ All of the above
- ✗ None of the above

11 The supervised learning problems can be grouped as ...

- ✗ Regression problems
- ✗ Classification problems
- ✓ All of the above
- ✗ None of the above

12 The unsupervised learning problems can be grouped as ...

- ✗ Clustering
- ✗ Association
- ✓ All of the above
- ✗ None of the above

MCQ (4/1)

13 Overfitting is a type of modelling error which results in the failure to predict future observations effectively or fit additional data in the existing model. Yes/No?

- × Probably
- ✓ Yes
- × No
- × Can not say

14 ... is the scenario when the model fails to decipher the underlying trend in the input data.

- ✓ Underfitting
- × Overfitting
- × All of the above
- × None of the above

15 The categories in which Machine learning approaches can be traditionally categorized are ...

- × Supervised learning
- × Unsupervised learning
- × Reinforcement learning
- ✓ All of the above

16 In general, to have a well-defined learning problem, we must identity which of the following

- × The class of tasks
- × The measure of performance to be improved
- × The source of experience
- ✓ All of the above

MCQ (5/1)

17 The average positive difference between computed and desired outcome values
- ✗ Root Mean Squared Error
- ✗ Mean Squared Error
- ✗ Mean Absolute Error
- ✓ Mean Positive Error

18 ... is used as an input to the machine learning model for training and prediction purposes.
- ✗ Target variable
- ✓ Feature vector
- ✗ All of the above
- ✗ None of the above

19 Simple regression assumes a ... relationship between the input attribute and output attribute.
- ✓ linear
- ✗ quadratic
- ✗ reciprocal
- ✗ inverse

20 The correlation between the number of years an emplyee has worked for a company and the salary of the employee is 0.75. What can be said about employee salary and years worked?
- ✗ There is no relationship between salary and years worked.
- ✓ Individuals that have worked for the company the longest have higher salaries.
- ✗ Individuals that have worked for the company the longest have lower salaries.
- ✗ The majority of employees have been with the company a long time.

MCQ (6/1)

31 Which machine learning models are trained to make a series of decisions based on the rewards and feedback they receive for their actions?

- ✗ Supervised learning
- ✗ Unsupervised learning
- ✓ Reinforcement learning
- ✗ All of the above

32 Which of the following is not a type of supervised learning?

- ✗ Classification
- ✗ Regression
- ✓ Clustering
- ✗ None of the above

33 As the amount of training data increases

- ✗ Training error usually increases and generalization error usually increases
- ✓ Training error usually increases and generalization error usually decreases
- ✗ Training error usually decreases and generalization error usually decreases
- ✗ Training error usually decreases and generalization error usually increases

34 Which of the following are not classification tasks?

- ✗ Find the gender of a person by analyzing his writing style
- ✗ Detect Pneumonia from Chest X-ray images
- ✓ Predict the price of a house based on floor area, number of rooms, etc.
- ✗ Predict whether there will be abnormally heavy rainfall next year

MCQ (7/1)

25 Which of the following is a categorical feature?

- ✗ Height of a person
- ✗ Price of petroleum
- ✗ Amount of rainfall in a day
- ✓ Mother tongue of a person

26 What is the use of validation dataset in Machine Learning?

- ✗ To train the machine learning model.
- ✓ To tune the hyperparameters of the machine learning model
- ✗ To evaluate the performance of the machine learning model
- ✗ None of the above

27 Which of the following criteria is typically used for optimizing in linear regression.

- ✗ Maximize the number of points it touches.
- ✗ Minimize the number of points it touches.
- ✓ Minimize the squared distance from the points.
- ✗ Minimize the maximum distance of a point from a line.

28 For two runs of K-Mean clustering, is it expected to get same clustering results?

- ✗ Yes
- ✓ No

MCQ (8/1)

㉙ Logistic Regression is used for ...

- ✕ regression purposes
- ✓ classification purposes
- ✕ all of the above
- ✕ none of the above

㉚ Which of the following methods do we use to best fit the data in Logistic Regression?

- ✕ Least Squared Error
- ✓ Maximum Likelihood
- ✕ Jaccard distance

㉛ When there is noise in data, which of the following options would improve the performance of the k-NN algorithm?

- ✓ Increase the value of k
- ✕ Decrease the value of k
- ✕ Changing value of k will not change the effect of the noise
- ✕ None of these

Some Useful Links

1. https://setosa.io/ev/
2. https://karpathy.ai/
3. http://yann.lecun.com/
4. https://www.hackingnote.com/
5. https://machinelearningmastery.com/
6. https://stanford.edu/ shervine/teaching/
7. https://www.ibm.com/downloads/cas/GB8ZMQZ3
8. https://colah.github.io/posts/2014-03-NN-Manifolds-Topology/

Temporary page!

LATEX was unable to guess the total number of pages correctly. As there was some unprocessed data that should have been added to the final page this extra page has been added to receive it.

If you rerun the document (without altering it) this surplus page will go away, because LATEX now knows how many pages to expect for this document.